

Task 2 - Dimensionality Reduction Methods

Ednalyn C. De Dios

August 24, 2023

Environment

- Python: 3.9.9
- Jupyter: 7.0.2

Part I - Research Question

A1. Propose one question relevant to a real-world organizational situation that you will answer by using principal component analysis (PCA).

Can the principal components of the customer base be identified using Principal Component Analysis?

A2. Define one goal of the data analysis. Ensure that your goal is reasonable within the scope of the scenario and is represented in the available data.

The ultimate goal of this data analysis is to reduce operating costs by increasing the efficiency of the organization's marketing efforts. We will use Principal Component Analysis (PCA) as a dimensionality reduction technique to identify the principal components of the data set. This will inform the decisions of stakeholders in matters where customer retention is involved, for example. Knowing which principal components accounts for the variance will provide the organization with advanced insight towards that customer's characteristics or behavior. Thereby increasing the effectiveness of marketing campaigns.

Part II - Method Justification

B1. Explain how PCA analyzes the selected data set. Include expected outcomes.

PCA analyzes the selected by first standardizing the data set and fitting it. Then, the variances of each component were plotted to discern the optimal number of k clusters using the elbow method. After choosing the number of k clusters, we then fit and transform the data set again with designated k number of clusters. The variances for each of the principal components were noted and then added to calculate the total variance captured by the principal components.

B2. Summarize one assumption of PCA.

PCA or Principal Component Analysis makes multiple assumptions about the data set but the biggest in my opinion is that variable should have correlation (Sharma, 2021). Data sets with high correlation between variables works best for PCA so that it can reduce the number of dimensions in the data set. Otherwise, there will be no reduction in dimensionality that can take place.

Part III - Data Preparation

C1. Identify the continuous dataset variables that you will need in order to answer the PCA question proposed in part A1.

```
In [1]: columns = ['Population',
                  'Children',
                  'Age',
                  'Income',
                  'Outage_sec_perweek',
                  'Email',
                  'Contacts',
                  'Yearly_equip_failure',
                  'Tenure',
                  'MonthlyCharge',
                  'Bandwidth_GB_Year',
                  ]
```

```
In [2]: # setting the random seed for reproducibility
import random
random.seed(493)

# for manipulating dataframes
import pandas as pd
import numpy as np

# for visualizations
```

```
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="whitegrid")
from IPython.display import Image

# for modeling
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

# to print out all the outputs of the cell
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# set display options
import warnings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
pd.set_option('display.max_colwidth', None)
```

```
In [3]: # read the csv file
df = pd.read_csv('churn_clean.csv')
df.info()
df.head()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 50 columns):

#	Column	Non-Null	Count	Dtype
0	CaseOrder	10000	non-null	int64
1	Customer_id	10000	non-null	object
2	Interaction	10000	non-null	object
3	UID	10000	non-null	object
4	City	10000	non-null	object
5	State	10000	non-null	object
6	County	10000	non-null	object
7	Zip	10000	non-null	int64
8	Lat	10000	non-null	float64
9	Lng	10000	non-null	float64
10	Population	10000	non-null	int64
11	Area	10000	non-null	object
12	TimeZone	10000	non-null	object
13	Job	10000	non-null	object
14	Children	10000	non-null	int64
15	Age	10000	non-null	int64
16	Income	10000	non-null	float64
17	Marital	10000	non-null	object
18	Gender	10000	non-null	object
19	Churn	10000	non-null	object
20	Outage_sec_perweek	10000	non-null	float64
21	Email	10000	non-null	int64
22	Contacts	10000	non-null	int64
23	Yearly_equip_failure	10000	non-null	int64
24	Techie	10000	non-null	object
25	Contract	10000	non-null	object
26	Port_modem	10000	non-null	object
27	Tablet	10000	non-null	object
28	InternetService	7871	non-null	object
29	Phone	10000	non-null	object
30	Multiple	10000	non-null	object
31	OnlineSecurity	10000	non-null	object
32	OnlineBackup	10000	non-null	object
33	DeviceProtection	10000	non-null	object
34	TechSupport	10000	non-null	object
35	StreamingTV	10000	non-null	object
36	StreamingMovies	10000	non-null	object
37	PaperlessBilling	10000	non-null	object
38	PaymentMethod	10000	non-null	object
39	Tenure	10000	non-null	float64
40	MonthlyCharge	10000	non-null	float64
41	Bandwidth_GB_Year	10000	non-null	float64
42	Item1	10000	non-null	int64
43	Item2	10000	non-null	int64
44	Item3	10000	non-null	int64
45	Item4	10000	non-null	int64
46	Item5	10000	non-null	int64
47	Item6	10000	non-null	int64
48	Item7	10000	non-null	int64
49	Item8	10000	non-null	int64

dtypes: float64(7), int64(16), object(27)
memory usage: 3.8+ MB

Out[3]:

	CaseOrder	Customer_id	Interaction	UID	City	S
0	1	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	e885b299883d4f9fb18e39c75155d990	Point Baker	
1	2	S120509	fb76459f-c047-4a9d-8af9-e0f7d4ac2524	f2de8bef964785f41a2959829830fb8a	West Branch	
2	3	K191035	344d114c-3736-4be5-98f7-c72c281e2d35	f1784cfa9f6d92ae816197eb175d3c71	Yamhill	
3	4	D90850	abfa2b40-2d43-4994-b15a-989b8c79e311	dc8a365077241bb5cd5ccd305136b05e	Del Mar	
4	5	K662701	68a861fd-0d20-4e51-a587-8a90407ee574	aabb64a116e83fdc4befc1fbab1663f9	Needville	

```
In [4]: for col in columns:
        # show outliers
        plt.boxplot(df[col])
        plt.title(col)
        fig = plt.figure(figsize =(10, 7))
```

Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a546f77760>, <matplotlib.lines.Line2D at 0x1a546f77a30>], 'caps': [<matplotlib.lines.Line2D at 0x1a546f77cd0>, <matplotlib.lines.Line2D at 0x1a546f77f70>], 'boxes': [<matplotlib.lines.Line2D at 0x1a546f774c0>], 'medians': [<matplotlib.lines.Line2D at 0x1a546f97250>], 'fliers': [<matplotlib.lines.Line2D at 0x1a546f974f0>], 'means': []}

Out[4]: Text(0.5, 1.0, 'Population')

Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a546fe43a0>, <matplotlib.lines.Line2D at 0x1a546fe4640>], 'caps': [<matplotlib.lines.Line2D at 0x1a546fe48e0>, <matplotlib.lines.Line2D at 0x1a546fe4b80>], 'boxes': [<matplotlib.lines.Line2D at 0x1a546fe4100>], 'medians': [<matplotlib.lines.Line2D at 0x1a546fe4e20>], 'fliers': [<matplotlib.lines.Line2D at 0x1a546ff4100>], 'means': []}

Out[4]: Text(0.5, 1.0, 'Children')

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a547038100>,
<matplotlib.lines.Line2D at 0x1a5470383a0>],
'caps': [<matplotlib.lines.Line2D at 0x1a546fd3040>,
<matplotlib.lines.Line2D at 0x1a547038370>],
'boxes': [<matplotlib.lines.Line2D at 0x1a547027e20>],
'medians': [<matplotlib.lines.Line2D at 0x1a547038700>],
'fliers': [<matplotlib.lines.Line2D at 0x1a5470389a0>],
'means': []}
```

```
Out[4]: Text(0.5, 1.0, 'Age')
```

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a547080910>,
<matplotlib.lines.Line2D at 0x1a547080a60>],
'caps': [<matplotlib.lines.Line2D at 0x1a547080d30>,
<matplotlib.lines.Line2D at 0x1a547080fd0>],
'boxes': [<matplotlib.lines.Line2D at 0x1a547080670>],
'medians': [<matplotlib.lines.Line2D at 0x1a54708d2b0>],
'fliers': [<matplotlib.lines.Line2D at 0x1a54708d550>],
'means': []}
```

```
Out[4]: Text(0.5, 1.0, 'Income')
```

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a5470d53d0>,
<matplotlib.lines.Line2D at 0x1a5470d5550>],
'caps': [<matplotlib.lines.Line2D at 0x1a5470d57f0>,
<matplotlib.lines.Line2D at 0x1a5470d5a90>],
'boxes': [<matplotlib.lines.Line2D at 0x1a5470d5130>],
'medians': [<matplotlib.lines.Line2D at 0x1a5470d5d30>],
'fliers': [<matplotlib.lines.Line2D at 0x1a5470d5fd0>],
'means': []}
```

```
Out[4]: Text(0.5, 1.0, 'Outage_sec_perweek')
```

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a5470f5be0>,
<matplotlib.lines.Line2D at 0x1a547119e50>],
'caps': [<matplotlib.lines.Line2D at 0x1a54712d130>,
<matplotlib.lines.Line2D at 0x1a54712d3d0>],
'boxes': [<matplotlib.lines.Line2D at 0x1a547119c70>],
'medians': [<matplotlib.lines.Line2D at 0x1a54712d670>],
'fliers': [<matplotlib.lines.Line2D at 0x1a54712d910>],
'means': []}
```

```
Out[4]: Text(0.5, 1.0, 'Email')
```

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a547170730>,
<matplotlib.lines.Line2D at 0x1a5471709d0>],
'caps': [<matplotlib.lines.Line2D at 0x1a547170c70>,
<matplotlib.lines.Line2D at 0x1a547170f10>],
'boxes': [<matplotlib.lines.Line2D at 0x1a5471705b0>],
'medians': [<matplotlib.lines.Line2D at 0x1a54717f1f0>],
'fliers': [<matplotlib.lines.Line2D at 0x1a54717f490>],
'means': []}
```

```
Out[4]: Text(0.5, 1.0, 'Contacts')
```

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a5471c71c0>,
<matplotlib.lines.Line2D at 0x1a5471c7460>],
'caps': [<matplotlib.lines.Line2D at 0x1a5471c7700>,
<matplotlib.lines.Line2D at 0x1a5471c79a0>],
'boxes': [<matplotlib.lines.Line2D at 0x1a5471b8ee0>],
'medians': [<matplotlib.lines.Line2D at 0x1a5471c7c40>],
'fliers': [<matplotlib.lines.Line2D at 0x1a5471c7ee0>],
'means': []}
```

```
Out[4]: Text(0.5, 1.0, 'Yearly equip_failure')
```

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a54720cb20>,
<matplotlib.lines.Line2D at 0x1a54720cdc0>],
'caps': [<matplotlib.lines.Line2D at 0x1a54721e0a0>,
<matplotlib.lines.Line2D at 0x1a54721e340>],
'boxes': [<matplotlib.lines.Line2D at 0x1a54720c880>],
'medians': [<matplotlib.lines.Line2D at 0x1a54721e5e0>],
'fliers': [<matplotlib.lines.Line2D at 0x1a54721e880>],
'means': []}
```

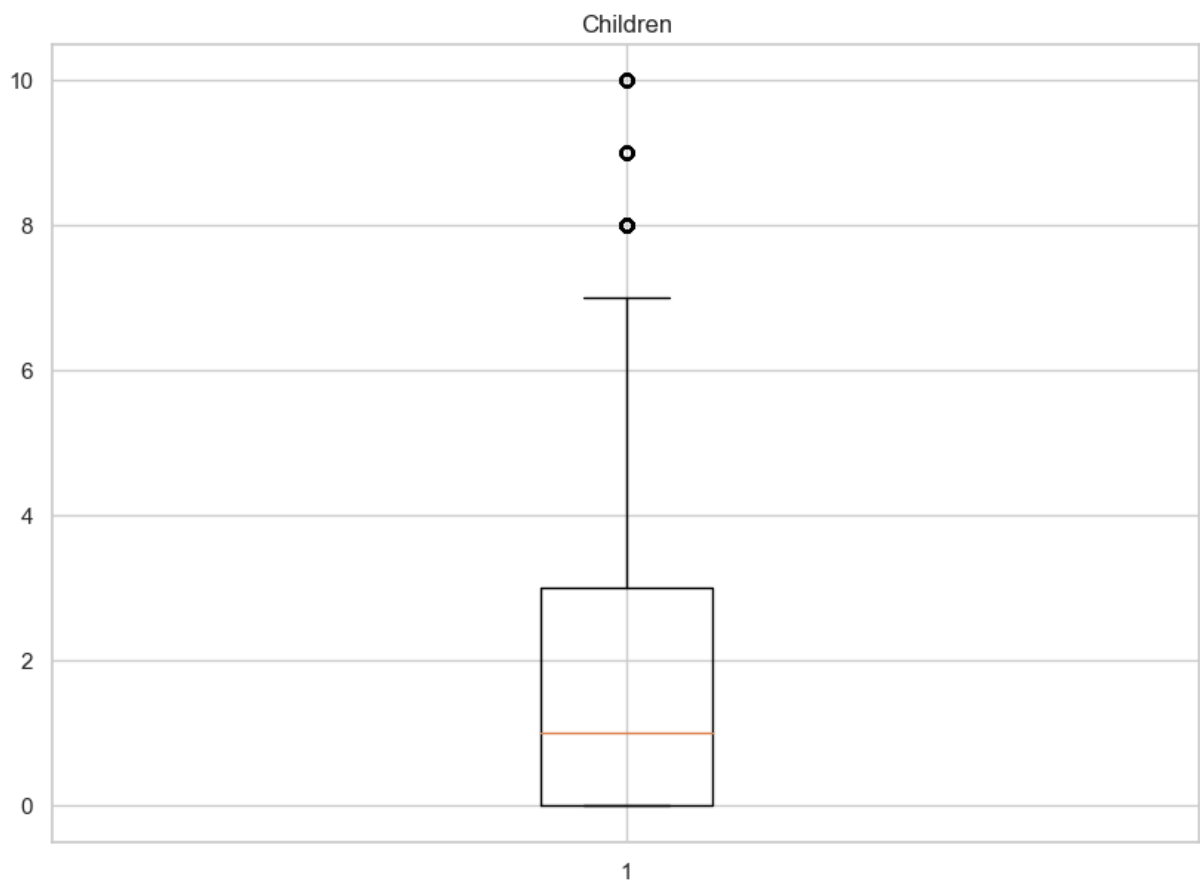
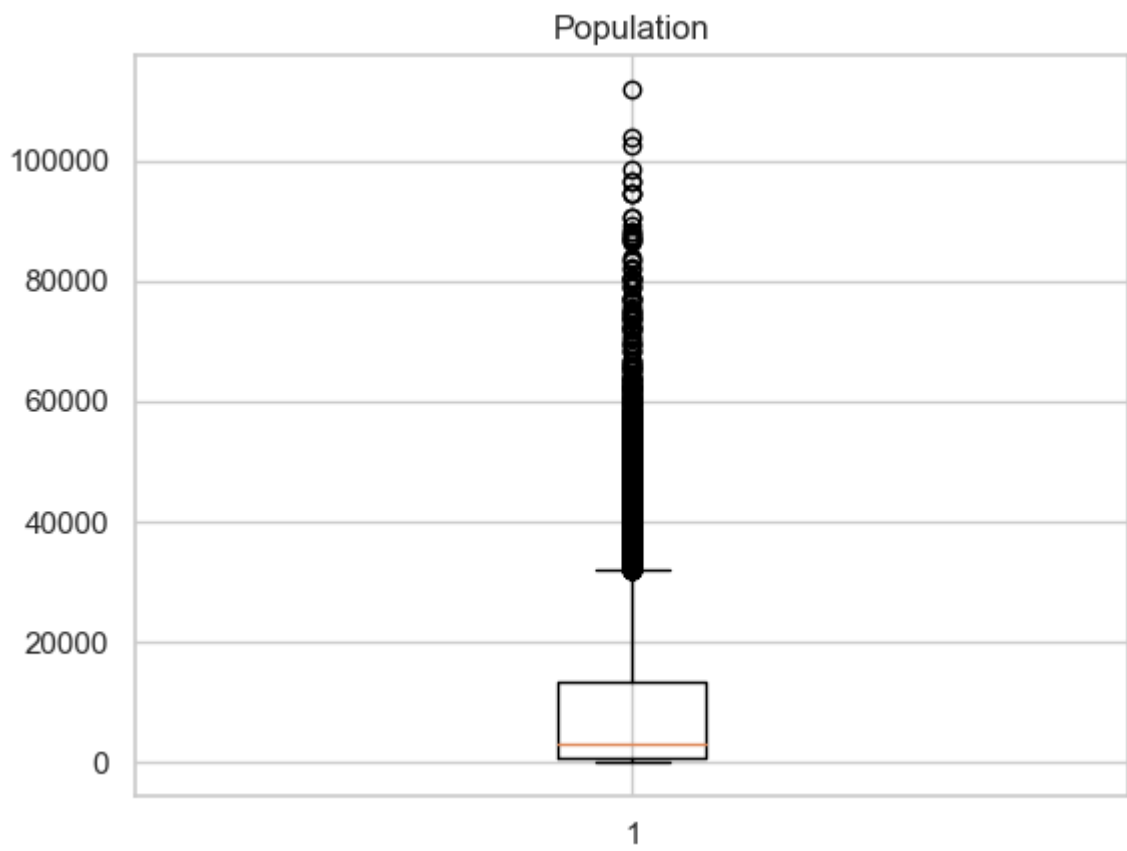
```
Out[4]: Text(0.5, 1.0, 'Tenure')
```



```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a5472636d0>,
<matplotlib.lines.Line2D at 0x1a547263970>],
'caps': [<matplotlib.lines.Line2D at 0x1a547263c10>,
<matplotlib.lines.Line2D at 0x1a547263eb0>],
'boxes': [<matplotlib.lines.Line2D at 0x1a547263430>],
'medians': [<matplotlib.lines.Line2D at 0x1a547273190>],
'fliers': [<matplotlib.lines.Line2D at 0x1a547273430>],
'means': []}
```


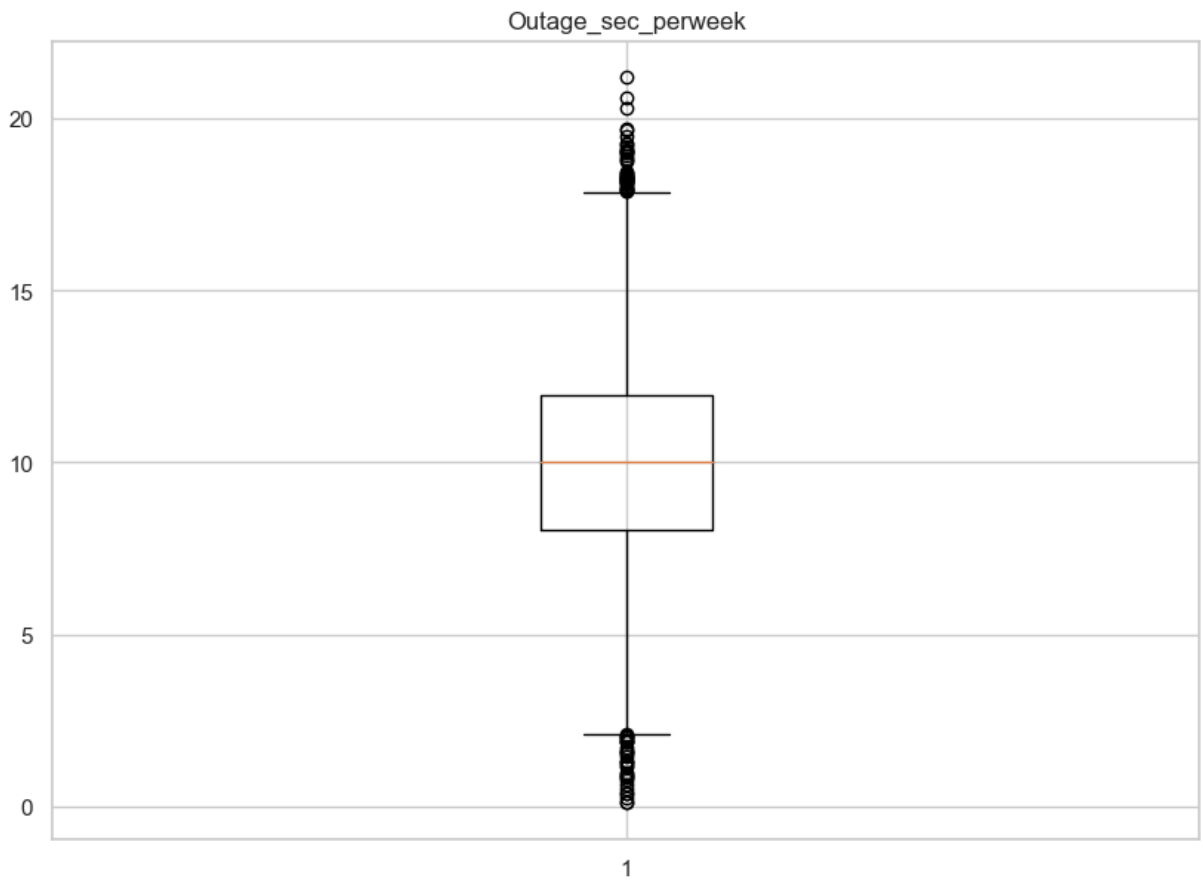
```
Out[4]: Text(0.5, 1.0, 'MonthlyCharge')
```

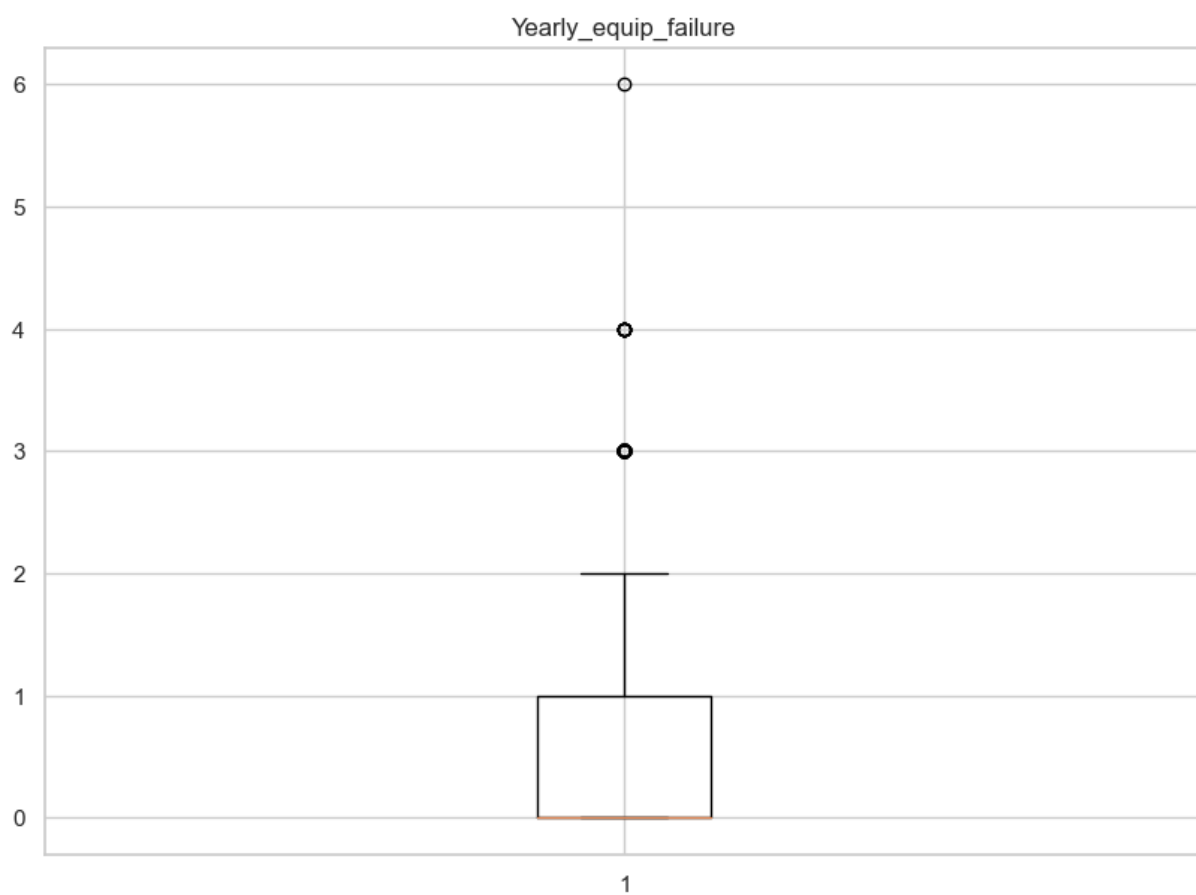

```
Out[4]: {'whiskers': [<matplotlib.lines.Line2D at 0x1a5475c9160>,
<matplotlib.lines.Line2D at 0x1a5475c9400>],
'caps': [<matplotlib.lines.Line2D at 0x1a5475c96a0>,
<matplotlib.lines.Line2D at 0x1a5475c9940>],
'boxes': [<matplotlib.lines.Line2D at 0x1a5472a8e80>],
'medians': [<matplotlib.lines.Line2D at 0x1a5475c9be0>],
'fliers': [<matplotlib.lines.Line2D at 0x1a5475c9e80>],
'means': []}
```



```
Out[4]: Text(0.5, 1.0, 'Bandwidth_GB_Year')
```














<Figure size 1000x700 with 0 Axes>

```
In [5]: # remove outliers in "Population"
df = df[df['Population'] <= 30000]
df.shape
```

Out[5]: (8957, 50)

```
In [6]: # remove outliers in "Children"
df = df[df['Children'] < 8]
df.shape
```

Out[6]: (8600, 50)

```
In [7]: # remove outliers in "Income"
df = df[df['Income'] <= 100000]
df.shape
```

Out[7]: (8255, 50)

```
In [8]: # remove outliers in "Outage_sec_perweek"
df = df[df['Outage_sec_perweek'] <= 17.5]
df.shape
```

Out[8]: (8198, 50)

```
In [9]: # remove outliers in "Email"
df = df[df['Email'] <= 20]
```

```
df.shape
```

```
Out[9]: (8187, 50)
```

```
In [10]: # remove outliers in "Contacts"
df = df[df['Contacts'] <= 5]
df.shape
```

```
Out[10]: (8181, 50)
```

```
In [11]: # remove outliers in "Yearly equip_failure"
df = df[df['Yearly equip_failure'] <= 2]
df.shape
```

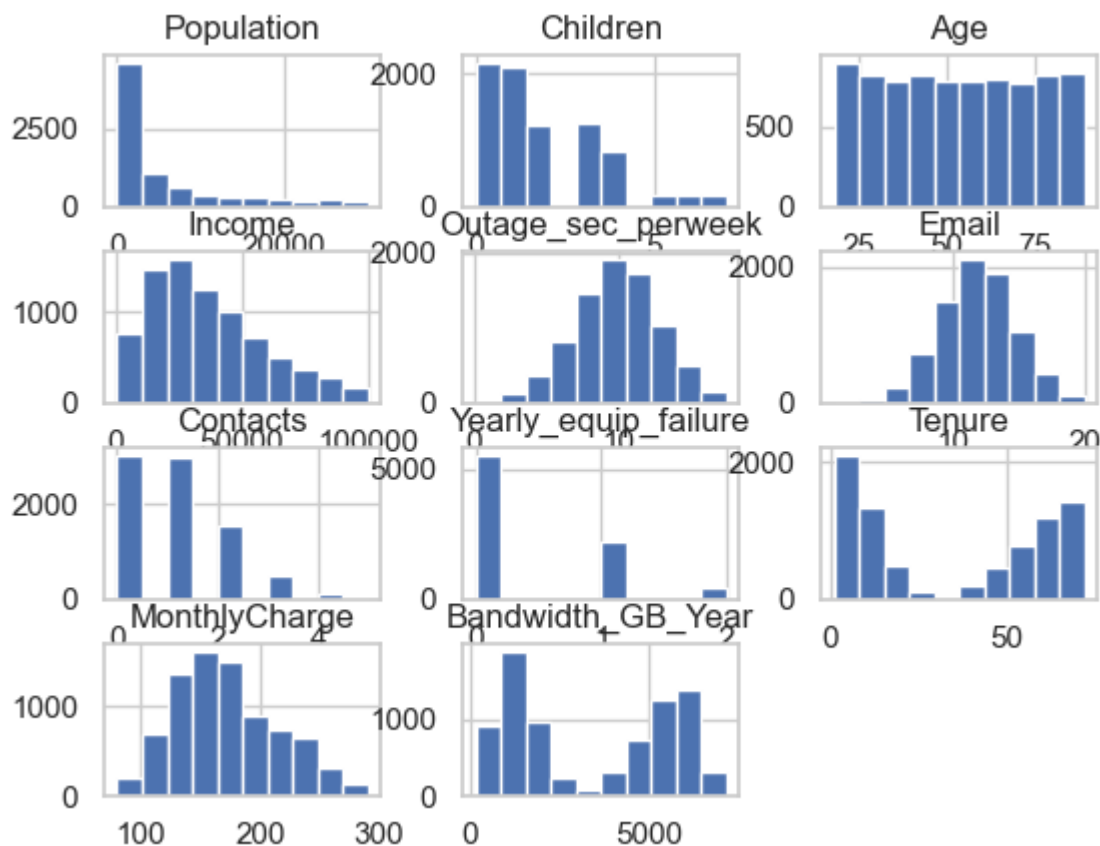
```
Out[11]: (8105, 50)
```

C2. Standardize the continuous dataset variables identified in part C1. Include a copy of the cleaned dataset.

```
In [12]: df = df[columns]
```

```
In [13]: # make histograms and save the plot
df[df.columns].hist()
```

```
Out[13]: array([[<Axes: title={'center': 'Population'}>,
  <Axes: title={'center': 'Children'}>,
  <Axes: title={'center': 'Age'}>],
  [<Axes: title={'center': 'Income'}>,
  <Axes: title={'center': 'Outage_sec_perweek'}>,
  <Axes: title={'center': 'Email'}>],
  [<Axes: title={'center': 'Contacts'}>,
  <Axes: title={'center': 'Yearly equip_failure'}>,
  <Axes: title={'center': 'Tenure'}>],
  [<Axes: title={'center': 'MonthlyCharge'}>,
  <Axes: title={'center': 'Bandwidth_GB_Year'}>, <Axes: >]],
  dtype=object)
```



```
In [14]: # scale the data
scaler = StandardScaler()

# apply scaler() to all the continuous column
scaled = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)

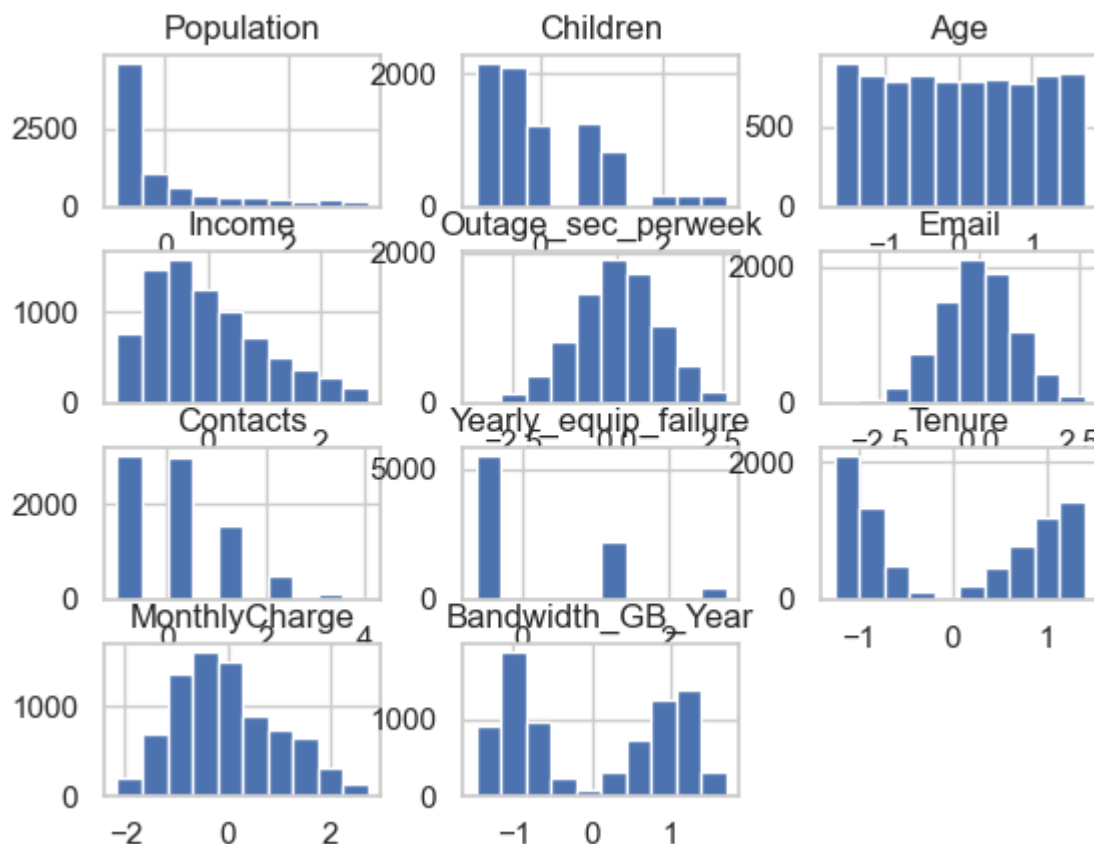
scaled.head()
```

```
Out[14]:
```

	Population	Children	Age	Income	Outage_sec_perweek	Email	Contacts
0	-0.763524	-1.069540	0.718481	-0.350998	-0.677041	-0.656913	-1.009708
1	0.640643	-0.478753	-1.266577	-0.657046	0.606316	0.007418	-1.009708
2	-0.264753	1.293609	-0.153008	-1.196872	0.279927	-0.989078	-1.009708
3	1.101638	-0.478753	-0.249840	-0.781101	1.715042	1.003913	1.024516
4	0.762874	-1.069540	1.444722	0.162809	-0.618718	1.336079	1.024516

```
In [15]: # make histograms and save the plot
scaled[scaled.columns].hist()
```

```
Out[15]: array([[<Axes: title={'center': 'Population'}>,
<Axes: title={'center': 'Children'}>,
<Axes: title={'center': 'Age'}>],
[<Axes: title={'center': 'Income'}>,
<Axes: title={'center': 'Outage_sec_perweek'}>,
<Axes: title={'center': 'Email'}>],
[<Axes: title={'center': 'Contacts'}>,
<Axes: title={'center': 'Yearly equip_failure'}>,
<Axes: title={'center': 'Tenure'}>],
[<Axes: title={'center': 'MonthlyCharge'}>,
<Axes: title={'center': 'Bandwidth_GB_Year'}>, <Axes: >]],
dtype=object)
```



```
In [16]: # save the prepared data set
scaled.to_csv('churn_prepared2.csv', index=False)
```

Part IV. Analysis

D1. Determine the matrix of all the principal components.

```
In [17]: pca = PCA(n_components=6, random_state=493)
model = pca.fit_transform(scaled)
explained_ratio = pca.explained_variance_ratio_
explained_ratio
```



```
Out[17]: array([0.1813204 , 0.09673   , 0.09385725, 0.0920565 , 0.09105452,
                0.09085826])
```

```
In [18]: print(pca.explained_variance_ratio_.cumsum())
```

```
[0.1813204  0.2780504  0.37190765 0.46396415 0.55501867 0.64587693]
```

```
In [19]: matrix = pd.DataFrame(model, columns=['PC1', 'PC2', 'PC3', 'PC4', 'PC5', 'PC6'])
```

```
In [20]: matrix.head()
matrix.tail()
```

```
Out[20]:
```

	PC1	PC2	PC3	PC4	PC5	PC6
0	-1.519209	-0.540523	-0.473303	0.861094	-0.652050	-1.097087
1	-1.653816	0.260500	1.735917	-0.171267	1.103416	-1.350004
2	-0.877558	-1.236562	0.909397	0.782003	-0.135945	-1.523216
3	-0.926511	1.736037	0.776245	-1.107354	-1.012047	0.374659
4	-1.930431	1.365848	-0.985916	0.598114	-0.916003	-0.457744

```
Out[20]:
```

	PC1	PC2	PC3	PC4	PC5	PC6
8100	1.217256	1.689024	-0.249704	-1.031767	-1.481937	0.314613
8101	0.825883	-0.672417	-1.407049	-1.078539	-0.437635	0.631326
8102	1.866298	-0.797156	0.315489	-0.836525	-0.198837	1.478551
8103	0.569396	-0.993756	-1.129926	-0.385718	0.151184	0.082555
8104	1.554721	1.639923	1.075307	-1.916243	0.736876	-0.692976

D2. Identify the total number of principal components using the elbow rule or the Kaiser criterion. Include a screenshot of the scree plot.

```
In [21]: plt.figure(figsize=(10, 8))
plt.plot(pca.explained_variance_ratio_)
plt.title('Scree Plot with Elbow')
plt.xlabel('Principal Components')
plt.ylabel('Explained Variance')
```


```
Out[21]: <Figure size 1000x800 with 0 Axes>
```

```
Out[21]: [<matplotlib.lines.Line2D at 0x1a5480d4550>]
```

```
Out[21]: Text(0.5, 1.0, 'Scree Plot with Elbow')
```

```
Out[21]: Text(0.5, 0, 'Principal Components')
```

```
Out[21]: Text(0, 0.5, 'Explained Variance')
```



Based on the scree plot above, two principal components are significant.

D3. Identify the variance of each of the principal components identified in part D2.

```
In [22]: def scree_plot(pca):
    num_components = len(pca.explained_variance_ratio_)
    ind = np.arange(num_components)
    vals = pca.explained_variance_ratio_

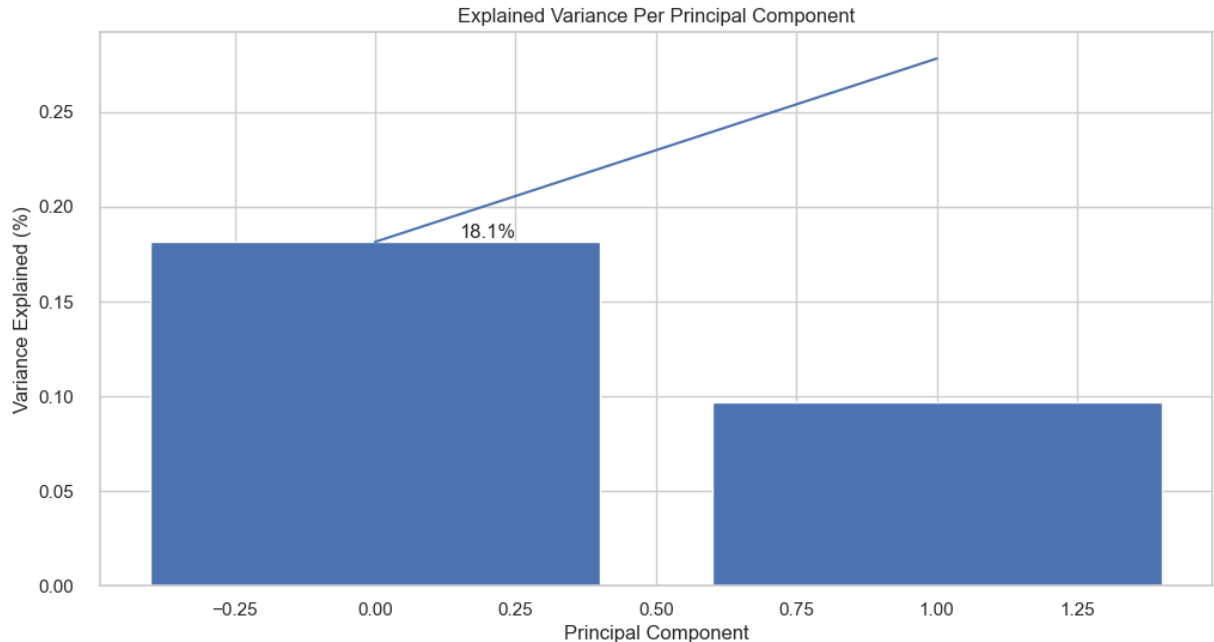
    plt.figure(figsize=(12, 6))
    ax = plt.subplot(111)
    cumvals = np.cumsum(vals)
    ax.bar(ind, vals)
    ax.plot(ind, cumvals)
    for i in range(num_components):
        if(i%10==0):
            ax.annotate(r"%s%%" % ((str(vals[i]*100)[:4])), (ind[i]+0.2, vals[i]), va="

    ax.xaxis.set_tick_params(width=0)
    ax.yaxis.set_tick_params(width=2, length=12)

    ax.set_xlabel("Principal Component")
```

```
ax.set_ylabel("Variance Explained (%)")
plt.title('Explained Variance Per Principal Component')
```

```
In [23]: # Re-apply PCA to the data while selecting for number of components to retain.
pca = PCA(n_components=2)
model = pca.fit_transform(scaled)
scree_plot(pca)
```



```
In [24]: print('Explained variation per principal component: {}'.format(pca.explained_variance_ratio_))
Explained variation per principal component: [0.1813204 0.09673 ]
```

D4. Identify the total variance captured by the principal components identified in part D2.

```
In [25]: print(str(round(sum(pca.explained_variance_ratio_)*100, 2)) + "%")
27.81%
```

D5. Summarize the results of your data analysis.

While we were able to identify the principal components of the data set, I believe that the result of the analysis is inconclusive. The total variance captured by the principal components identified in part D2 is only 27.81%. This is simply not enough for any actionable insight to base on.

Part V. Attachments

E. Record the web sources used to acquire data or segments of third-party code to support the analysis. Ensure the web sources are reliable.

- <https://github.com/ecdedios/code-snippets/blob/main/notebooks/master.ipynb>
- [https://github.com/ecdedios/d212-data-mining-ii/blob/main/D212%20Performance%20Assessment%20Task%201%20\(Rav.%200\).ipynb](https://github.com/ecdedios/d212-data-mining-ii/blob/main/D212%20Performance%20Assessment%20Task%201%20(Rav.%200).ipynb)
- <https://github.com/microbhai/ChurnAnalysis/blob/master/PrincipalComponentAnalysis.ipynb>
- <https://www.datacamp.com/tutorial/principal-component-analysis-in-python>

F. Acknowledge sources, using in-text citations and references, for content that is quoted, paraphrased, or summarized.

- <https://medium.com/data-science-on-customer-churn-data/pca-or-principal-component-analysis-on-customer-churn-data-d18ca60397ed>